

# Homework 02

## Keith Wampler

AI Tool Used: Gemini 2.5 Pro

### Task 1 - NRI Data Cleaning

1. Import the NRI data. Ensure that the FIPS code variable ('STCOFIPS') is correctly identified as a string / character variable. Otherwise, the leading zeros will be removed.

```
In [1]: import pandas as pd
import numpy as np

# path to the NRI file
nri_path = '/Users/keith/Documents/code/Intro to ML 2025/data/raw/NRI_Table_

# import the data w/ FIPS column as a string type
nri_df = pd.read_csv(nri_path, dtype={'STCOFIPS': str})

# check
print(nri_df.head())
```

	OID_	NRI_ID	STATE	STATEABBRV	STATEFIPS	COUNTY	COUNTYTYPE	\
0	1	C01001	Alabama	AL	1	Autauga	County	
1	2	C01003	Alabama	AL	1	Baldwin	County	
2	3	C01005	Alabama	AL	1	Barbour	County	
3	4	C01007	Alabama	AL	1	Bibb	County	
4	5	C01009	Alabama	AL	1	Blount	County	

  

	COUNTYFIPS	STCOFIPS	POPULATION	...	WNTW_EALS	WNTW_EALR	\
0	1	01001	58764	...	15.784587	Very Low	
1	3	01003	231365	...	56.205509	Relatively Moderate	
2	5	01005	25160	...	18.632002	Relatively Low	
3	7	01007	22239	...	13.308573	Very Low	
4	9	01009	58992	...	23.645930	Relatively Low	

  

	WNTW_ALRB	WNTW_ALRP	WNTW_ALRA	WNTW_ALR_NPCTL	WNTW_RISKV	\
0	2.687716e-07	7.410082e-09	8.725777e-06	10.461158	8494.906508	
1	1.268231e-09	2.287120e-08	1.548360e-07	13.339523	65619.701638	
2	5.788050e-07	2.347236e-08	7.606598e-07	16.125039	15501.730335	
3	9.014679e-07	1.270300e-08	1.202015e-05	16.991643	7496.186940	
4	5.268425e-07	1.482016e-08	2.002965e-07	12.039616	17175.160729	

  

	WNTW_RISKS	WNTW_RISKR	NRI_VER
0	12.217626	Very Low	March 2023
1	52.083996	Relatively Low	March 2023
2	19.535476	Very Low	March 2023
3	11.104041	Very Low	March 2023
4	21.444480	Very Low	March 2023

[5 rows x 465 columns]

2. Subset the NRI data to include only the 5-digit state/county FIPS code and all columns ending with '\_AFREQ' and '\_RISKR'. Each of these columns represents a different hazard type.

```
In [2]: # find columns ending with '_AFREQ' or '_RISKR'
selected_cols = ['STCOFIPS'] + [col for col in nri_df.columns if col.endswith('_AFREQ') or col.endswith('_RISKR')]

# create the subset
nri_subset = nri_df[selected_cols]

# check
print(nri_subset.head())
```

	STCOFIPS	AVLN_AFREQ	AVLN_RISK	CFLD_AFREQ	CFLD_RISK	\
0	01001	NaN	Not Applicable	NaN	Not Applicable	
1	01003	NaN	Not Applicable	3.684142	Relatively Low	
2	01005	NaN	Not Applicable	NaN	Not Applicable	
3	01007	NaN	Not Applicable	NaN	Not Applicable	
4	01009	NaN	Not Applicable	NaN	Not Applicable	

  

	CWAV_AFREQ	CWAV_RISK	DRGT_AFREQ	DRGT_RISK	ERQK_AFREQ	...
0	0.0	No Rating	25.969774	Relatively Low	0.000431	...
1	0.0	No Rating	12.353442	Relatively Moderate	0.000338	...
2	0.0	No Rating	43.956953	Relatively Low	0.000227	...
3	0.0	No Rating	28.894501	Very Low	0.000790	...
4	0.0	No Rating	28.152598	Relatively Low	0.000817	...

  

	TRND_AFREQ	TRND_RISK	TSUN_AFREQ	TSUN_RISK	VLCN_AFREQ
0	0.480184	Relatively Moderate	NaN	Not Applicable	NaN
1	0.953140	Relatively Moderate	NaN	Insufficient Data	NaN
2	0.739018	Relatively Moderate	NaN	Not Applicable	NaN
3	0.586160	Relatively Moderate	NaN	Not Applicable	NaN
4	0.710332	Relatively Moderate	NaN	Not Applicable	NaN

  

	VLCN_RISK	WFIR_AFREQ	WFIR_RISK	WNTW_AFREQ	WNTW_RISK
0	Not Applicable	0.000035	Very Low	0.433437	Very Low
1	Not Applicable	0.002229	Relatively Moderate	0.182759	Relatively Low
2	Not Applicable	0.000038	Very Low	0.185759	Very Low
3	Not Applicable	0.000040	Very Low	0.743034	Very Low
4	Not Applicable	0.000035	Very Low	0.866873	Very Low

[5 rows x 37 columns]

3. Create a table / dataframe that, for each hazard type, shows the number of missing values in the '\_AFREQ' and '\_RISK' columns.

```
In [3]: # extract unique hazard prefixes
hazards = sorted(list(set([col.split('_')[0] for col in nri_subset.columns ]))

# Create a dictionary to store missing value counts
missing_counts = {}

for hazard in hazards:
    afreq_col = f"{hazard}_AFREQ"
    riskr_col = f"{hazard}_RISK"
    missing_counts[hazard] = {
        'AFREQ_missing': nri_subset[afreq_col].isnull().sum(),
        'RISKR_missing': nri_subset[riskr_col].isnull().sum()
    }

# convert the dictionary to a df printing
missing_df = pd.DataFrame.from_dict(missing_counts, orient='index')

print(missing_df)
```

	AFREQ_missing	RISKR_missing
AVLN	3023	0
CFLD	2646	0
CWAV	0	0
DRGT	7	0
ERQK	0	0
HAIL	7	0
HRCN	918	0
HWAV	0	0
ISTM	229	0
LNDS	40	0
LTNG	123	0
RFLD	0	0
SWND	7	0
TRND	7	0
TSUN	3103	0
VLCN	3125	0
WFIR	88	0
WNTW	0	0

4. Create a new column in the original data table indicating whether or not 'AVLN\_AFREQ' is missing or observed. Show the cross-tabulation of the 'AVLN\_AFREQ' missingness and 'AVLN\_RISKR' columns (including missing values). What do you observe?

```
In [4]: # create a new column that is True if AVLN_AFREQ is null, and False otherwise
nri_df['AVLN_AFREQ_missing'] = nri_df['AVLN_AFREQ'].isnull()

# make cross-tabulation table
# dropna=False ensures we see counts of missing values in the RISKR column as well
cross_tab = pd.crosstab(
    nri_df['AVLN_AFREQ_missing'],
    nri_df['AVLN_RISKR'],
    dropna=False
)

print(cross_tab)
```

AVLN_RISKR	Not Applicable	Relatively High	Relatively Low	\
AVLN_AFREQ_missing				
False	0	15	52	
True	3023	0	0	

  

AVLN_RISKR	Relatively Moderate	Very High	Very Low
AVLN_AFREQ_missing			
False	33	9	99
True	0	0	0

Observation: AVLN\_AFREQ is missing in every single case where AVLN\_RISKR is "Not Applicable". This suggests that when an avalanche risk isn't applicable to a county (e.g., Florida), its frequency isn't recorded.

5. Assuming that a risk that is “not applicable” to a county has an annualized frequency of 0, impute the relevant missing values in the ‘\_AFREQ’ columns with 0.

```
In [5]: # make list of all AFREQ columns
afreq_columns = [col for col in nri_df.columns if col.endswith('_AFREQ')]

# fill missing values in these columns with 0
nri_df[afreq_columns] = nri_df[afreq_columns].fillna(0)

# check
print(f"missing values in AVLN_AFREQ after imputation: {nri_df['AVLN_AFREQ']}
```

missing values in AVLN\_AFREQ after imputation: 0

## Task 2 - SVI Data Cleaning

1. Import the SVI data. Ensure that the FIPS code is correctly identified as a string / character variable. Otherwise, the leading zeros will be removed.

1. Subset the SVI data to include only the following columns: ST, STATE, ST\_ABBR, STCNTY, COUNTY, FIPS, LOCATION, AREA\_SQMI, E\_TOTPOP, EP\_POV150, EP\_UNEMP, EP\_HBURD, EP\_NOHSDP, EP\_UNINSUR, EP\_AGE65, EP\_AGE17, EP\_DISABL, EP\_SNGPNT, EP\_LIMENG, EP\_MINRTY, EP\_MUNIT, EP\_MOBILE, EP\_CROWD, EP\_NOVEH, EP\_GROUPQ, EP\_NOINT, EP\_AFAM, EP\_HISP, EP\_ASIAN, EP\_AIAN, EP\_NHPI, EP\_TWOMORE, EP\_OTHERRACE

```
In [6]: # path to your SVI file
svi_path = '/Users/keith/Documents/code/Intro to ML 2025/data/raw/SVI_2022_L

# import the SVI data w/ FIPS is a string
svi_df = pd.read_csv(svi_path, dtype={'FIPS': str})

# columns to keep
svi_columns_to_keep = [
    'ST', 'STATE', 'ST_ABBR', 'STCNTY', 'COUNTY', 'FIPS', 'LOCATION', 'AREA_
    'E_TOTPOP', 'EP_POV150', 'EP_UNEMP', 'EP_HBURD', 'EP_NOHSDP', 'EP_UNINSU
    'EP_AGE65', 'EP_AGE17', 'EP_DISABL', 'EP_SNGPNT', 'EP_LIMENG', 'EP_MINRT
    'EP_MUNIT', 'EP_MOBILE', 'EP_CROWD', 'EP_NOVEH', 'EP_GROUPQ', 'EP_NOINT'
    'EP_AFAM', 'EP_HISP', 'EP_ASIAN', 'EP_AIAN', 'EP_NHPI', 'EP_TWOMORE', 'E
]

# make subset
svi_subset = svi_df[svi_columns_to_keep]

# check
print(svi_subset.head())
```

	ST	STATE	ST_ABBR	STCNTY	COUNTY	FIPS	\
0	1	Alabama	AL	1001	Autauga County	01001	
1	1	Alabama	AL	1003	Baldwin County	01003	
2	1	Alabama	AL	1005	Barbour County	01005	
3	1	Alabama	AL	1007	Bibb County	01007	
4	1	Alabama	AL	1009	Blount County	01009	

  

	LOCATION	AREA_SQMI	E_TOTPOP	EP_POV150	...	EP_NOVEH
\						
0	Autauga County, Alabama	594.454786	58761	20.2	...	4.0
1	Baldwin County, Alabama	1589.861817	233420	18.3	...	2.3
2	Barbour County, Alabama	885.007619	24877	37.7	...	11.7
3	Bibb County, Alabama	622.469286	22251	29.0	...	7.5
4	Blount County, Alabama	644.890376	59077	22.9	...	4.8

  

	EP_GROUPQ	EP_NOINT	EP_AFAM	EP_HISP	EP_ASIAN	EP_AIAN	EP_NHPI	\
0	0.9	10.9	19.6	3.2	1.1	0.1	0.0	
1	1.5	10.9	8.3	4.8	0.9	0.2	0.0	
2	12.0	31.8	46.9	4.8	0.5	0.3	0.0	
3	6.4	20.2	20.7	2.9	0.3	0.1	0.0	
4	1.0	16.9	1.2	9.7	0.2	0.1	0.2	

  

	EP_TWOMORE	EP_OTHERRACE
0	3.3	0.2
1	3.1	0.4
2	1.8	1.2
3	1.7	0.1
4	2.8	0.1

[5 rows x 33 columns]

2. Create a table / dataframe that shows the number of missing values in each column. (Hint: if you wrote a function for Task 1, you can reuse it here.)

```
In [7]: # sum of nulls for each column
missing_svi = svi_subset.isnull().sum()

# check
print(missing_svi)
```

```

ST          0
STATE       0
ST_ABBR     0
STCNTY      0
COUNTY     0
FIPS        0
LOCATION     0
AREA_SQMI   0
E_TOTPOP    0
EP_POV150   0
EP_UNEMP     0
EP_HBURD    0
EP_NOHSDP   0
EP_UNINSUR   0
EP_AGE65    0
EP_AGE17    0
EP_DISABL   0
EP_SNGPNT   0
EP_LIMENG   0
EP_MINRTY   0
EP_MUNIT    0
EP_MOBILE   0
EP_CROWD     0
EP_NOVEH    0
EP_GROUPQ   0
EP_NOINT    0
EP_AFAM     0
EP_HISP     0
EP_ASIAN    0
EP_AIAN     0
EP_NHPI     0
EP_TWOMORE  0
EP_OTHERRACE 0
dtype: int64

```

### Task 3 - Data Merging

1. Identify any FIPS codes that are present in the NRI data but not in the SVI data and vice versa. Describe any discrepancies and possible causes? What to these discrepancies, if any, mean for interpreting results based on the merged dataset moving forward?

```

In [8]: # get the set of FIPS codes from each df
nri_fips = set(nri_df['STCOFIPS'])
svi_fips = set(svi_df['FIPS'])

# find FIPS codes in NRI but not in SVI
nri_only = nri_fips - svi_fips
print(f"FIPS codes in NRI only: {nri_only}")
print(f"Count: {len(nri_only)}\n")

# find FIPS codes in SVI but not in NRI
svi_only = svi_fips - nri_fips

```

```
print(f"FIPS codes in SVI only: {svi_only}")
print(f"Count: {len(svi_only)}")
```

FIPS codes in NRI only: {'72111', '72143', '78030', '72131', '72043', '72147', '72087', '72031', '72141', '72065', '72035', '72049', '09013', '72097', '69120', '72139', '69110', '72013', '72115', '72119', '72095', '72051', '60010', '09003', '72007', '72091', '72023', '60050', '09005', '72027', '72113', '72081', '09011', '72015', '72089', '72083', '78020', '09015', '72011', '72135', '72075', '72117', '72009', '72061', '72071', '72099', '72133', '72077', '72107', '72037', '72101', '78010', '09007', '72017', '72103', '72021', '72145', '72039', '72073', '72129', '72127', '72093', '72121', '72079', '72105', '72109', '72033', '72019', '72137', '09009', '72057', '72054', '72067', '72025', '72151', '66010', '72125', '72153', '72047', '72123', '72059', '60020', '69100', '72055', '72085', '09001', '72029', '72063', '72069', '72053', '72149', '72001', '72005', '72003', '72041', '72045'}

Count: 96

FIPS codes in SVI only: {'09120', '09160', '09140', '09180', '09110', '09170', '09150', '09190', '09130'}

Count: 9

The two datasets don't perfectly align because the SVI data includes Puerto Rico while the NRI does not, and the NRI also uses a few outdated county codes. When these files are merged, you get incomplete rows for mismatched counties, and data for specific locations like military bases is simply included within their surrounding county's record. This means your results could be misleading unless you first harmonize the county codes and account for these geographic differences in your analysis.

2. Merge the NRI and SVI data on the FIPS code. Use an outer join to keep all counties in the final dataset.

```
In [9]: # rename columns before merging
nri_cleaned = nri_df.rename(columns={'STCOFIPS': 'FIPS'})
svi_cleaned = svi_subset # Already subsetted in Task 2

# outer merge
merged_df = pd.merge(
    nri_cleaned,
    svi_cleaned,
    on='FIPS',
    how='outer'
)

# show size of merged df
print(merged_df.shape)
```

(3240, 498)

3. Create a table / dataframe that shows the number of missing values in each column of the merged dataset.

```
In [10]: # get missing value counts for the merged df
print(merged_df.isnull().sum())
```



```

OID_          9
NRI_ID        9
STATE_x       9
STATEABBRV    9
STATEFIPS     9
..
EP_ASIAN      96
EP_AIAN       96
EP_NHPI       96
EP_TWOMORE    96
EP_OTHERRACE  96
Length: 498, dtype: int64

```

## Task 4 - Data Analysis

1. For each numerical variable in the merged dataset, plot a histogram showing the distribution of values. (Hint: write a function to make the histogram for a single variable, then use a loop or apply function to make the histograms for all numerical variables.)

```

In [11]: import matplotlib.pyplot as plt
import seaborn as sns

# select only numerical columns for plotting
numerical_cols = merged_df.select_dtypes(include=np.number).columns

# figure to hold the plots
plt.figure(figsize=(20, 200))

for i, col in enumerate(numerical_cols):
    plt.subplot(len(numerical_cols) // 4 + 1, 4, i + 1)
    sns.histplot(merged_df[col], kde=False, bins=50)
    plt.title(col)
    plt.xlabel('')
    plt.ylabel('')

plt.tight_layout()
plt.show()

```













